

Characterizing and Curating Conversation Threads: Expansion, Focus, Volume, Re-entry

Lars Backstrom
Facebook
lars@fb.com

Jon Kleinberg
Cornell University
kleinber@cs.cornell.edu

Lillian Lee
Cornell University
llee@cs.cornell.edu

Cristian Danescu-Niculescu-Mizil
Stanford/Max Planck Institute SWS
cristiand@cs.stanford.edu

ABSTRACT

Discussion threads form a central part of the experience on many Web sites, including social networking sites such as Facebook and Google Plus and knowledge creation sites such as Wikipedia. To help users manage the challenge of allocating their attention among the discussions that are relevant to them, there has been a growing need for the algorithmic curation of on-line conversations — the development of automated methods to select a subset of discussions to present to a user.

Here we consider two key sub-problems inherent in conversational curation: length prediction — predicting the number of comments a discussion thread will receive — and the novel task of re-entry prediction — predicting whether a user who has participated in a thread will later contribute another comment to it. The first of these sub-problems arises in estimating how interesting a thread is, in the sense of generating a lot of conversation; the second can help determine whether users should be kept notified of the progress of a thread to which they have already contributed. We develop and evaluate a range of approaches for these tasks, based on an analysis of the network structure and arrival pattern among the participants, as well as a novel dichotomy in the structure of long threads. We find that for both tasks, learning-based approaches using these sources of information yield improvements for all the performance metrics we used.

Categories and Subject Descriptors: H.2.8: Data Mining

General Terms: Measurement; Experimentation; Theory

Keywords: user-generated content, comment threads, threads, Facebook, Wikipedia, conversations, likes, feed ranking, recommendation, on-line communities, social networks, discussions

1. INTRODUCTION

Many Web sites are organized around a continuously evolving set of discussion threads. This style of interaction is a key component of on-line groups and message boards, social networking sites such as Facebook and Google Plus, and the workflow of collaborative projects such as Wikipedia and open-source development. In all these cases, a user must continuously decide how to allocate his

or her attention to a range of relevant discussions, and this can be a challenging task when content arrives at a rapid rate.

A growing number of sites are helping users address this challenge through the algorithmic *curation* of discussion threads, automatically selecting which threads to bring to a user's attention at any given point in time. A canonical example is Facebook's News Feed — for users with a sufficient number of active friends on the site, an unfiltered stream of all stories generated by friends is generally much less valuable to the user than a ranked and filtered version of the stream that attempts to highlight the stories estimated to be most engaging to the user.

The problem of curating discussion threads is thus a wide-ranging one in the context of applications, but it is one which for the most part has not been systematized in prior research. Our goal in this paper is to facilitate such a systematization, by identifying and formalizing two important sub-problems in conversational curation, and then developing and evaluating techniques to address them. For our evaluation, we use discussion threads from two sites where such threads form a core part of the experience: discussions among users on Facebook and discussions among editors on Wikipedia. As on many other sites, threads on Facebook and Wikipedia can be conceptualized as an initial *post* and a subsequent sequence of *comments*; we will use this terminology in what follows.

The present work: Two problems in conversational curation.

We now describe the two problems that we study, together with their motivation as components of conversational curation.

1. *Length prediction*: given the initial portion of a thread (a post and the first few comments following it), how well can we predict the eventual length of the thread? We use this length prediction problem as a concretely formulated proxy for the general issue of estimating the level of interest a thread will generate, based on observation of its early stages.
2. *Re-entry prediction*: given the initial portion of a thread and the identity of one of the commenters, how well can we predict whether this commenter will contribute another comment later in the thread? This is a key issue in determining whether to keep a user notified of the progress on a thread once he or she has contributed to it — some threads have the structure of a conversation where users are motivated to return repeatedly, while others involve each user contributing once (for example, to offer congratulations or condolences) but then not returning.

Taken together, these two problems cover a set of central issues in conversational curation: identifying threads that will generate sustained interest, so as to be able to highlight them to users, and

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

WSDM'13, February 4–8, 2013, Rome, Italy.

Copyright 2013 ACM 978-1-4503-1869-3/13/02 ...\$15.00.

recognizing whether a thread is something that a contributing user will want to continue to follow as it evolves.

We develop techniques for these problems by first analyzing the structure of threads, and then formulating a set of properties that we in turn use for the prediction tasks.

We begin by investigating the following issue, on data described in §2. Intuitively, one feels from experience that there are two distinct types of long threads: those that become long because a small group of people engage in an extensive conversation via the comments, and those that become long because many users each contribute a single comment. A canonical example of the latter would begin with a post in which a user announces a major life event, and then many friends contribute congratulations in the comments as in a wedding guestbook. We refer to the first type of thread as *focused*, and the second type as *expansionary*.

But is this notion of two types simply one’s perception of two extremes of a broad distribution, or is there quantitative evidence for it? We find (§3) in fact that threads genuinely exhibit this two-type effect: for long threads, the distribution of the number of distinct commenters is bimodal, with threads either dominated by a very small number of distinct users, or by a sequence of commenters who generally do not return to the thread after commenting once. In addition to providing what is, to our knowledge, the first evidence for this basic dichotomy, this finding helps reinforce the importance of our second problem — re-entry prediction — by establishing that active discussion threads can vary considerably in the extent to which participants are interested in returning after their initial contribution.

In order to build a framework for approaching our two basic problems, we begin by studying (§4) a range of related thread properties. One of the most useful of these is the thread’s *arrival pattern* — the ordering by which new entrants into the thread are interleaved with returning participants. Formalizing this notion allows us to work with relaxed versions of the two extremes of focused and expansionary threads discussed above, and to explore the region that interpolates between them. We also study network and temporal structure: whether the first few commenters are linked within a social network, and how quickly after the post do they arrive in real-time; both convey information about the future trajectory of the thread.

We incorporate these properties into a machine learning approach for predicting length (§5) and re-entry (§6). Evaluating the prediction performance enables us to identify the features that are most effective for our two problems. At a high level, we find that the structure of the arrival pattern is the most useful for re-entry prediction, while temporal properties together with the arrival pattern give the strongest performance for length prediction.

Next, in §7, we explore a probabilistic model of participant re-entry related to the dichotomy between focused threads and expansionary ones. Clearly some styles of post tend to lead to one type of thread or the other, but for other kinds of posts, one sees both types of threads emerge; for example, the same shared link to a news story can generate a focused thread when it is shared among one set of users and an expansionary thread when it is shared among a different set. It is therefore natural to ask whether a type of symmetry-breaking can arise directly from the dynamics of a discussion itself — that is, whether there is a simple probabilistic generative model capable of producing both focused and expansionary threads over different realizations of its random trajectory. We show how to construct such a model from plausible assumptions about turn-taking and new entrants in discussion threads; the model exposes interesting connections between discussion threads and nonlinear urn processes.

In §8, we review related work on the dynamics of on-line discussions. For now, we note that the general issue of thread length has been studied, using different techniques, in contexts distinct from ours — primarily for comments on blog and news sites, where essentially all threads are expansionary, with many participants who typically contribute only once or very few times each [14, 23, 25, 26]. In contrast, our approach incorporating the notion that there can be multiple structurally distinct types of long threads is suited to settings where the participants maintain long-running relationships with one another. These structural distinctions also provide a core part of the motivation for re-entry prediction, which is a key issue for organizing conversations in these settings; the problem of re-entry prediction has not, to our knowledge, been formulated or studied previously.

2. DATA AND BASIC DEFINITIONS

We use data from Facebook and Wikipedia to construct three distinct populations of users whose discussion threads we study. We choose Facebook as perhaps the most well-known example of a post-plus-comments interface for socially-oriented conversations. Conversations among Wikipedia editors form a contrasting case that has also received research attention [7, 13, 17]: the discussions are task-oriented, as opposed to socially-oriented, and there is no formal structure imposed on conversations by the interface; nonetheless, they can still be naturally treated as instances of comment threads.

For completeness, we briefly describe the structure of these discussion threads at a general level. On Facebook, we study instances in which a user posts a status update, and then other users with permission to comment on the status update contribute comments to it. On Wikipedia, editors interact on *talk-pages* to discuss issues concerning articles, projects or Wikipedia policies. Each editor has the option of hosting a talk-page, and most active users do.

On both Facebook and Wikipedia, we will refer to the status update or initiating text as the *post*; the sequence of comments that follows the post will be called the *comment thread*, and the post together with all the comments will be called the *full thread*. The poster together with the commenters in a full thread will be called the full thread’s *participants*. The number of items in the thread (including the post in the case of a full thread, but not in a comment thread) will be called its *length* or its *volume*; we use these two terms synonymously.

From Facebook, we first selected 100,000 users uniformly at random from the population of US Facebook users. We will refer to this set \mathcal{U} in our analysis as the *uniform* Facebook population. Also, out of all US Facebook users who posted between 200 and 300 status updates over an 80-day period, we randomly selected 100,000 of these heavily engaged users. We will refer to this set \mathcal{A} as the *high-activity* Facebook population. For both \mathcal{U} and \mathcal{A} , we study the comment threads associated with all their posts during the same 80-day period. All Facebook data was used anonymously, and all analysis was done in aggregate.

Our Wikipedia data is derived from the corpus of Danescu-Niculescu-Mizil et al. [7]. We used 118,447 conversation threads of length at least 1 (to discard posts made by automated bots, which never attract responses) which took place asynchronously on the talk-pages hosted by 6,555 highly active editors; posts average 2.12 comments. We also use the content of the talk-pages to assess the existence of an interactional link between a given pair of Wikipedia editors: we say that two editors are linked if at least one of them added a post or comment on the other’s editor talk page. Our Wikipedia data will be available at

http://www.mpi-sws.org/~cristian/Echoes_of_power.html.




Figure 1: Heat map (best viewed in color) of the density functions on distinct commenters, uniform Facebook population. For each thread prefix length k , there is a peak in density (lighter color) at a small number and a second peak approximately at the maximum number of distinct participants d .

3. FOCUSED VS. EXPANSIONARY THREADS

If a post leads to a long comment thread, then it is one that attracts a great deal of attention and so is likely of interest; thus, the thread-length prediction problem is crucial to the curating of conversations. In thinking about how to bring long threads to users’ attention, though, a natural question is whether there are sub-classes of such conversations that should be treated differently.

This question leads us to conjecture that there is a dichotomy between *expansionary* high-activity threads, created by the one-time actions of many different “drive-by” commenters, versus *focused* high-activity threads, reflecting a high-level of repeated engagement among relatively few people. In this section, we provide supporting evidence for this conjecture and discuss its consequences.

Distinct Participants: Two Local Maxima. To investigate the validity of our conjecture, we consider how the number of distinct participants in a thread is distributed. To do so, we must account for two issues. First, we do not want the idiosyncratic actions of any one high-volume user to dominate the quantities involved, so we work with a macro-averaged function.¹ Second, the possible number of distinct participants in a thread depends on the thread’s length, and so we need to parametrize by it.

Thus, formally, for a population of users \mathcal{P} , let \mathcal{P}_k be the set of users who authored at least one post having comment thread length at least k . For each user $u \in \mathcal{P}_k$, we take all full threads associated with a post by u that produced at least k comments, and we truncate each of these threads to the prefix consisting of just the post and the first k comments. Let $\delta_u(k)$ be the average number of distinct participants in all these prefixes of full threads initiated by u . (For a given such prefix, the number can range from 1 — the original poster contributed all of the comments as well — to $k + 1$ — all commenters are distinct, and the original poster didn’t comment.) We then define $\Delta_k^*(d)$ to be the fraction of users $u \in \mathcal{P}_k$ for whom $\lfloor \delta_u(k) \rfloor = d$. Note that Δ_k^* is a density function. In what follows, for brevity we will sometimes refer to it simply as an average or an expectation, with the understanding that this refers in fact to a macro-averaged quantity.

In these terms, our conjecture can be expressed as follows: for threads of sufficient length k , the density function $\Delta_k^*(d)$ should

have (at least) two local maxima: one at a small value of d , i.e., $d \ll k$, and one at a large value of d , i.e., $d \approx k + 1$.

In Figure 1, we show the family of density functions Δ_k^* for $k \in 1, 2, \dots, 50$ on our uniform Facebook population \mathcal{U} . The densities are drawn as a heat map, with column k representing the density function Δ_k^* . We see that as k increases, $\Delta_k^*(d)$ is first maximized at $d = 4$, reflecting the dominant role of the focused effect; but then, a second local maximum emerges at a value of d very close to k . For the population \mathcal{A} of high-activity Facebook users, we see essentially the same effect, including the two local maxima at $d = 4$ and d close to k (figure omitted for space).

Although the smaller data volume makes it more difficult to discern the effect on Wikipedia, when we group together the possible values of d into contiguous intervals, we find significant evidence of two local maxima there too. To quantify the effect on Wikipedia, we compare quantiles of Δ_k^* , defining $f_k(p, q) = \sum_{pk \leq d \leq qk} \Delta_k^*(d)$.

We find that as k increases (in particular, considering $k \geq 15$), we have $f_k(0, \frac{1}{4}) > f_k(\frac{1}{4}, \frac{1}{2})$ and $f_k(\frac{3}{4}, 1) > f_k(\frac{1}{2}, \frac{3}{4})$. This inequality is consistent with Figure 1, where the density function is larger at the two extremes than in comparably-sized intervals in between.

Consequences: Predict Both Length and Re-entry. As argued earlier, conversational-curation systems should contain a thread-length prediction component. But our new observation about the distinction between expansionary and focused threads shows that long threads can differ significantly in the extent to which a commenter will want to return to contribute a second time. This heterogeneity in long threads motivates the formulation of our second task, *re-entry prediction*: determining whether a given participant in the thread is likely to contribute again. To our knowledge, re-entry in on-line conversations is a problem that has not been previously formalized or studied.

Length and re-entry are important, and distinct, issues in the task of conversational curation. Length prediction, since it provides information about the amount of attention a thread is likely to receive, helps in assessing whether a user should be made aware of the thread at all. Re-entry prediction, on the other hand, provides information about how to keep a user informed of the evolution of the thread once he or she has already contributed to it: a high re-entry probability indicates that the user may well want to know about subsequent comments, so that he or she can contribute in response to them.

Predicting a particular user’s re-entry is different from predicting whether the thread itself will be focused or expansionary. While very few users re-enter an expansionary thread by definition, it is easily possible for a user u to contribute to a thread that later becomes dominated by a back-and-forth discussion among a small set of other participants; in this case, the thread is focused, but user u ’s re-entry probability might be low. Predicting re-entry provides a concrete recommendation with respect to a given user, in a way that predicting whether a thread will be focused or expansionary does not.

We note that re-entry prediction is focused on a user’s *production* of comments — specifically, whether the user will write another comment in the future. An interesting open question is to consider the analogous prediction task for a user’s *consumption* of comments. In particular, a user might be interested in continuing to read comments on a thread as they come in, despite having no intention of contributing again. (Consider a string of congratulatory messages on a life event that include interesting side information, such as personal reminiscences or clever quips.)

¹The results turn out to be similar for the micro-averaged analog.

4. EARLY PARTICIPANTS: SOCIAL, SEQUENTIAL, AND TEMPORAL STRUCTURES

In this section, we show how properties of the initial participants in a thread can provide information about the thread’s later dynamics, thus laying the groundwork for the features in our subsequent prediction experiments. First (§4.1), we show that the presence or absence of social links among the initial participants in a thread turns out to provide useful information, though in interestingly different ways for different settings. Second (§4.2), inspired by our expansionary vs. focused analysis in §3, which introduces the importance of re-entry, we develop a novel representation for the sequence of participant contributions. Third (§4.3), we demonstrate that how fast the initial commenters arrive provides important information about the eventual number of comments, though its connection with re-entry probability is less clear.

4.1 Links Among the Initial Participants

We first consider comment threads in our Facebook High-Activity population (outcomes are analogous for the Uniform set), focusing on threads with at least two comments and where the first two commenters are distinct from each other and from the post’s author. The tension between the focused and the expansionary effects has a natural reflection in the relationship between these first two commenters. If they are friends, then interest in the post might be limited to a particular portion of the poster’s social neighborhood. That is, interest in such a post could have limited *reach*, which could restrict thread length. At the same time, though, there might also be increased potential for an extended conversation to ensue as friends interact, which would lead to a longer thread.

The top-left plot in Figure 2 shows that in fact, these Facebook threads are significantly longer when the first two commenters are friends.

We can further validate this hypothesized effect of conversational interaction by examining a related mechanism in which the role of interaction is much more limited. The Facebook “like” feature is very useful for this purpose. Users can respond to a post not just by commenting on it but also by clicking the *like* (thumbs-up) button, which provides a one-bit endorsement of the content. Thus, likes are a light-weight communication alternative to comments, and we can consider “like threads” — the sequence of likes arriving on a post — as the corresponding analog of comment threads. But there is a crucial difference: in like threads, there is no analog to the back-and-forth interaction that characterizes conversational interaction.

When the first two likers in a like thread are distinct, how is the eventual length of the like thread affected by whether these two users are friends? The top-middle plot of Figure 2(b) shows that, in the absence of repeated interactions to offset its consequences, a limited-reach effect is clear: the like thread is shorter when the first two likers are linked.²

Applying the same analysis to Facebook thread prefixes of length $k = 3, 4$, and 5 yields very similar results. For space reasons, we only depict the case $k = 3$ (bottom-left and bottom-middle plots in Figure 2), but the results are that expected length of comment threads continues to increase almost perfectly monotonically in the number of edges among the first k commenters when they are all distinct. Also, we find completely analogous results for re-entry in Facebook threads: the re-entry of the first participant increases

²We note that measuring the number of *distinct* commenters shows the same limited-reach effect: although the comment thread is longer when the first two commenters are linked, the total number of distinct commenters is smaller.




Figure 2: Thread length vs. number of connections between the first k commenters when they are distinct. More-connected commenters make for longer Facebook (FB) comment threads but shorter “like threads” and Wikipedia (WK) discussion threads. (Y-axes not aligned to heighten trend visibility.)

strongly with the number of edges among the first k commenters when they are all distinct.

What about Wikipedia comment threads? As depicted in the rightmost column of Figure 2, in this domain, it is *not* the case that more connections between the first participants leads to longer threads, although the available data here is quite sparse. We conjecture that the root cause for the striking contrast to Facebook may be the task-oriented nature of the setting, in which conversations may be less discursive, and editors who have interacted in the past may be more conversationally efficient in reaching a conclusion.

It is interesting to note that earlier work of Ugander et al. considered the level of connectedness among the set of users who appear in an invitation to join Facebook [24]; invitations that displayed users who were not linked to each other had higher overall conversion rates than invitations that displayed linked users. While invitations and comment threads are clearly different in nature, they both involve opportunities to engage a user in the activity of the site; whether there is a deeper relationship between the connectedness of commenters here and the connectedness of inviters in that setting is an interesting open question.

4.2 Arrival Patterns

Having looked at the *number* of distinct commenters, and at the *graph structure* on the first few participants in the case when they are distinct, we now develop a general method for representing the precise *sequence* of arrivals of the first few participants, and show that these sequences have the potential to be very useful features in our prediction tasks.

For a comment thread t , let t_i (for $i = 1, 2, \dots$) denote the identity of author of the i^{th} comment in the thread. We now define the following encoding $\gamma(t)$ of comment thread t . $\gamma(t)$ is a sequence of non-negative integers; the i^{th} entry in the sequence, denoted $\gamma(t)_i$, is equal to 0 if t_i is the author of the post that started the thread (returning to the thread in this case as the i^{th} commenter), and otherwise $\gamma(t)_i$ is equal to the value of j such that t_i is the j^{th} distinct commenter to take part in t . In what follows, we refer to $\gamma(t)_i$ as the *ID code* of commenter t_i . We will also use the term *arrival pattern* to refer generically to any prefix of $\gamma(t)$ (including the full sequence). Figure 3 illustrates these concepts via two sample discussions, exemplifying that focused threads should have arrival patterns in which some back-and-forth between two participants is evident, whereas expansionary threads should have arrival patterns in which all ID codes occur very few times, mostly just once.

Can early (i.e., short) arrival patterns serve as useful features for our prediction tasks? Before describing our full experiments (detailed in the next sections), it is useful to show some preliminary evidence of these patterns’ potential utility.

Facebook High-Activity, Length-5 arrival patterns		
pattern	1 re-enters	% of occ.
1,0,1,0,1	55.2%	19.2
1,0,1,0,0	47.5%	2.8
1,0,1,0,2	26.7%	4.9
1,0,1,2,0	26.1%	4.1
1,0,2,0,2	16.5%	4.6
1,2,0,2,0	14.6%	3.1
1,0,2,3,0	12.5%	1.9
1,2,0,3,0	11.5%	2.0
1,0,2,0,3	10.9%	2.6
1,2,3,4,5	5.6%	3.6
sum: 48.8		

Wikipedia, Length-5 arrival patterns		
pattern	1 re-enters	% of occ.
1,0,1,1,0	60.2%	2.4
1,1,0,1,0	58.6%	2.8
1,0,1,0,0	55.3%	4.8
1,0,0,1,0	52.2%	5.5
1,1,1,1,1	47.5%	2.3
1,0,1,2,1	46.0%	3.2
1,0,1,0,2	45.8%	2.7
1,2,1,2,1	41.1%	5.1
1,0,1,0,1	38.8%	27.0
1,2,3,4,5	7.2%	1.7
sum: 57.5		

Facebook High-Activity, length-9 arrival patterns		
pattern bins	1 re-enters	% of occ.
#0:3, #1:6	67.7%	1.7
#0:4, #1:5	66.9%	12.1
#0:5, #1:4	65.5%	4.0
#0:3, #1:4, #2:2	56.8%	2.1
#0:3, #1:3, #2:3	50.9%	1.7
#0:4, #1:4, #2:1	47.6%	5.2
#0:4, #1:3, #2:2	38.5%	3.5
#0:4, #1:3, #2:1, #3:1	28.2%	2.0
#0:4, #1:2, #2:3	22.3%	2.8
#0:4, #1:1, #2:4	9.6%	2.7
sum: 37.8		

Table 1: Left and middle: the most common length-5 arrival patterns on Facebook, accounting for 48.6% of the occurrences of all possible such arrival patterns, and on Wikipedia, accounting for 57.5% of all occurrences of all possible such patterns. The patterns are sorted by the percentage of corresponding threads in which the user with ID code 1 returns to the thread to comment again. “% of occ.”: percentage of threads of length ≥ 5 prefixed by that pattern. **Right:** The same for the most common length-9 arrival patterns, except that patterns have been binned by counts of ID codes since there are many possible length-9 patterns. For example, “#0:3, #1:6” = the set of length-9 patterns where ID code 0 occurs 3 times and ID code 1 occurs 6 times, in any order. (Some populations omitted for brevity or due to data sparseness.)

focused thread		expansionary thread
Mary:	Anyone there	James: we’re engaged!
Mary:	?	Dina: congrats!
Don:	me	Fred: congrats!
Pat:	not me	Mia: great!!!
Don:	v funny	Moe: great!
Pat:	i know	James: Thanks guys :)
...
Length-2 arrival pattern: 0,1		Length-2 arrival pattern: 1,2
Length-5 arrival pattern: 0,1,2,1,2		Length-5 arrival pattern: 1,2,3,4,0

Figure 3: Example conversations demonstrating our arrival-pattern coding scheme for the comment portion of threads.

First, we see whether different (early) arrival patterns tend to correspond to different thread lengths. Figure 4 shows, for each of our three populations, the (macro-averaged) length of threads whose length-two prefixes correspond to each of the five possible length-two patterns; the fact that the mean thread lengths fall in mostly disjoint confidence intervals indicates that the patterns do have predictive value.³

Second, we see whether different arrival patterns tend to correspond to differing re-entry probabilities, focusing on the chance that the user with ID code 1 (i.e., the first commenter who isn’t the original poster) subsequently re-joins the thread by adding another comment. Table 1 demonstrates that arrival patterns carry significant information about ID code 1’s re-entry probability. In all the populations shown, it appears that guestbook-style patterns containing many distinct ID codes tend to results in noticeably lower re-entry probabilities. For Facebook, we also see a strong positive correlation between the number of times ID code 1 appears in an arrival pattern and the likelihood that ID code 1 will subsequently appear again.

4.3 Timing effects

Our analysis thus far has considered the sequence of commenters without any information about the speed at which they arrive in real

³We note that the two most frequent arrival patterns in all three populations are (1,0) and (1,2), which is interesting because (1,0) corresponds to the canonical turn-taking structure in a pairwise conversation, while (1,2) is the canonical sequence of successive new arrivals — a further reflection of our focused/expansionary dichotomy.




Figure 4: In all populations, the 95%-confidence-intervals for mean thread length for the five possible length-2 arrival patterns — indicated as labels on the intervals — are almost all disjoint. Grey/dashed intervals indicate rare arrival patterns (at most 1% of threads), so the long interval involved in the single overlap (0,0 in Facebook Uniform) is for a sparse situation.

time. We now show some basic results establishing that this type of temporal structure contains important information about the length and re-entry properties of threads; in the next section, we use this information as part of our prediction methods.

In Figure 5, we see (black curve) that the longer it takes for the first comment to arrive on an initial post, the shorter the thread, presumably because “late” first comments correspond to less overall activity around the post. But note that timing isn’t everything: beyond a certain point, the probability that the first commenter re-enters a thread (green curve) becomes approximately independent of the first-comment arrival time lag.

5. PREDICTING THREAD LENGTH

We now engage in the two main prediction tasks of this paper. Recall that the first task, which we describe in this section, is to




Figure 5: Black (left axis, top curve): the longer it takes for the original post to attract its first comment, the lower the expected thread length. Green (right axis, bottom curve): in contrast, the probability that the first commenter re-enters the thread is eventually independent of the first comment’s arrival time. The $y = 3$ and $y = 20\%$ lines are included to allow for visual comparison of the real data curves with theoretical curves in which the arrival lag has no effect.

predict thread length, as an indication of how much interest a post will eventually generate, given the state of the thread at a certain early point. For example, we ask, given a thread initiated by a Facebook user posting a status update to their friends that has already accumulated 5 comments, how well can we predict the final length of the thread?⁴ The results of the second task, to predict whether a user that has already participated in a thread will later re-enter that same thread, are described in Section 6.

For thread-length prediction, we formulate two concrete tasks on the Wikipedia and Facebook High-Activity datasets (the results on the Facebook Uniform dataset are similar, but with larger error bars). From the set of all posts made by the active Facebook users, we selected the subset of posts that received at least 5 comments, and randomly reserved 50% of them for evaluation, using the other 50% for feature exploration. This gave us a test data set of 1,996,624 posts. Out of these, we chose a threshold of 8 comments to create an approximately balanced binary prediction problem: given the state of the thread after five posts, will the thread eventually receive at least 8 comments? (55.25% of the posts are in the positive class.) Similarly, for Wikipedia, we look at all talk-page posts that have received at least two comments, and ask, will they receive a third one? In this case, our data is smaller, with only 44,732 items in the test set (54.55% of which are in the positive class).

5.1 Features (used here and in §6)

The features we employed are summarized in Table 2. The first three sets are based on our discussion above of links between participants (§4.1), arrival patterns (§4.2), and timing effects (§4.3). We describe the other two sets now.

An important question is whether the textual features of the original post are more or less effective for this task than the non-textual features we have already described. To investigate this issue, we elected to gather a small, presumably general set of such “Original post terms” features via text regression, which has previously been employed for blog comment-volume prediction [26]. Specifically, we used J. M. White’s TextRegression R package, which employs linear regression with elastic-net regularization [9], run on a set of posts disjoint from the training and test data used for

⁴Naturally, for this task, we use only features that can be derived from the state of the thread when it had 5 comments.

LINKS	
$edges_prev[i]^*$	Number of links from commenter to previous commenters
$mutual_poster[i]^*$	Number of links from commenter to users linked to the original poster
ARRIVAL PATTERNS	
$id_code[i]$	commenter ID code as described in §4.2
$uniq_comm[i]$	Unique commenters through comment i
TIME	
$time[i]$	Time taken for the first i comments to arrive
TEXT REGRESSION FEATURES	
$Orig_post_terms$	“comment”, “agree”, etc.: see §5.1
MISC	
$num_words[i]$	Number of words in comment i
$num_chars[i]$	Number of characters in comment i
$question[i]^*$	Comment i has a ‘?’
$exclaim[i]^*$	Comment i has a ‘!’
$likes[i]^*$	Num likes on original post before comment i is made
$comment_likes[i]^*$	Num likes on comments before comment i

Table 2: Features used in our prediction experiments. For each indexed feature, we also build a comparable feature for the original post when it makes sense (the id_code for the original post is always 0 and so is omitted but, for example, the length of the original post in words or characters is meaningful). Features marked with * were applied only for Facebook data.

classification. 50 terms were selected for the Facebook data — among them were “comment” and “anybody” (positive coefficient for thread length), and “re-post” and URLs (negative coefficient). Among the 30 selected terms for Wikipedia were “agree” (positive) and “thank” (negative).

Also, preliminary pilot studies revealed a set of fairly intuitive miscellaneous features, listed in the last section of Table 2, that are potentially correlated with thread length. For instance, one might expect that on average, posts containing a question mark pose questions that prompt comments as responses.

5.2 Performance Results

Our testing methodology was: for a given set of features and train/test set, create bagged decision trees with 60 trees trained on independent samples of the training data; then, apply the bagged decision trees on the disjoint test set.

Our main method was to use all the features described in Table 2. We compared its performance against the following two **baselines**. The *positive-percentage bias baseline* chooses an item’s label randomly with bias equal to the percentage of test items in the positive class (55.52% in the Facebook case, 54.55% in the Wikipedia case). The *text-regression baseline* uses only the $Orig_post_terms$ features chosen via text regression as described in §5.1.

The performance of our method versus the two baselines is shown in Table 3. Clearly, the combined use of participant-link, arrival-pattern, timing, and other information yields the best results for all five of our performance metrics. The small set of text-regression features extracted from the original post sometimes did worse than the positive-percentage bias baseline.⁵

⁵This is consonant with De Choudhury et al. [8], who remark that

		ACC	AUC	RMSE	APR	CXE
FB	Pos.-% bias baseline	.552	.500	.497	.550	.992
	Text baseline	.537	.529	.503	.568	1.01
	All our features	.672	.729	.457	.758	.872
Wiki	Pos.-% bias baseline	.548	.500	.498	.549	.993
	Text baseline	.488	.505	.517	.550	1.06
	All our features	.595	.627	.486	.661	.958

Table 3: Main thread-length prediction results. Bold = best performance per dataset, under various metrics: ACC: accuracy (for FB active: after 5 comments, predicting whether the thread achieves length ≥ 8 ; for Wiki: after 2 comments, predicting whether an additional comment will occur). AUC: area under the ROC curve. RMSE: root mean square error. APR: mean average precision. CXE: cross-entropy.

Key Facebook Features. To better understand the individual factors contributing to the length of a comment thread, we perform stepwise forward feature selection. In iteration j of this algorithm, we create working feature set F_j by finding the best single feature to add to the set F_{j-1} to maximize our objective function, area under the ROC curve (AUC). Because it only selects a single feature at a time, this method prevents us from adding more than a single copy of highly correlated features, and the order that the features are installed gives us some insight into the nature of these comment threads.

Table 4 shows the features selected by this process for the Facebook dataset. There are three things worth noting from these results. The first is that a relatively small set of features contributes almost all of the predictive value. In particular, the amount of time it takes for the first five comments to arrive (the `TIME:time[5]` feature) is highly indicative of whether or not the thread will eventually reach 8 comments. Second, most of the key features come from the fifth comment. Thus, when predicting whether or not the thread will continue, one should focus on the most recent activity of the thread. We highlight this in Figure 6, where we show the prediction performance when using only the subsets of features derived from a single message in the thread, ranging from the original post ($x = 0$) to the fifth comment. The third item of note is the fact that the link-based features do not have much effect. We believe this is because they are low-recall, in the sense that we only showed in §4.1 that they are useful when all the early commenters in the thread are distinct.

Given the strength of the time feature, it is interesting to ask what would be the effect of its removal. The combination of the other features is unable to make up for the loss of temporal information: removing that key feature, the AUC drops from 0.729 to 0.588. With or without that feature, and even if we slice the data to predicting only for a fixed `TIME:time[5] $\in [15m, 20m)$` , the relative ordering of the other features remains more or less unchanged.

Key Wikipedia Features. In the case of Wikipedia, we see a somewhat similar ordering to the features. Again, the features regarding the most recent comment (here, the second one) are the most predictive of future comments. Here, however, we find that the *length* of the second comment is the most important feature, followed by the time to the second comment, and the ID code of the second and first commenters. Beyond these first four features, the rela-

“textual analyses ... alone are not adequate to capture conversational interestingness because [they] do not consider the dialogue structure between users”.

Feature added	AUC
<code>TIME:time[5]</code>	0.6954
<code>+ARRIVAL PATTERN:uniq_comm[5]</code>	0.7053
<code>+MISC:num_words[5]</code>	0.7138
<code>+TIME:time[3]</code>	0.7214
<code>+MISC:question[5]</code>	0.7256
<code>+ARRIVAL PATTERN:id_code[5]</code>	0.7258
<code>+ARRIVAL PATTERN:uniq_comm[4]</code>	0.7260

Table 4: Results of stepwise forward feature selection on Facebook. Each row represents performance for all features listed in that row and above.




Figure 6: Performance when predicting using only the features derived from a *single* comment (or original post for $x = 0$). The later the comment, the more informative.

tively small size of the dataset makes the predictive power of other features unclear.

6. PREDICTING THREAD RE-ENTRY

Here we examine our second, and novel, prediction task: given the initial portion of a thread and the identity of one of the commenters, how well can we predict whether that commenter will contribute another comment later in the thread? As noted earlier, the idea here is to determine whether to keep a user notified of the progress on a thread after they have commented on it already — there are some threads for which the user might want to actively return to the discussion. For space reasons, we can only provide an overview of results here, omitting detailed feature analysis.

For simplicity, we focus on the following two (related) questions. Recall that we use ID code 0 for the original poster of a thread, and ID code 1 for the first commenter other than the original poster, assuming there is such a commenter. (a) Assuming that ID code 1 occurs in the length-5 arrival-pattern prefix, does that user ever appear again? (The value 5 was used in our thread-length prediction problem as well.) (b) The same, but for the first 9 comments. We use the same features as in the previous task; see §5.1 and Table 2.

Using cross-validation, we find (Table 5) that the performance on the full feature set for Facebook is an AUC of 0.855 for the 9 comment version, and 0.808 for the 5 comment version of the task. Using the same feature selection methodology described above, we find that the most important features are the identities of the individuals posting the comments (`id_code[i]`), and especially the identities of the most recent few commenters. The time between the two most recent comments also plays an important role, as the longer it takes, the slower the conversation is moving, and the more likely it is to come to an end.

7. MODELING THREAD RE-ENTRY

Having gained some empirical understanding of thread re-entry, including relatively good performance at predicting it, we now seek

		AUC (x-val)
FB (after 5 comments)	Pos.-% bias baseline	.500
	Text baseline	.520
	Our features	.808
FB (after 9 comments)	Pos.-% bias baseline	.500
	Text baseline	.525
	Our features	.855
Wiki (after 5 comments)	Pos.-% bias baseline	.500
	Text baseline	.494
	Our features	.644

Table 5: Main thread-re-entry prediction cross-validation results. Bold marks the best performance per dataset.

to develop further theoretical understanding of re-entry by formulating a set of probabilistic generative models that produce arrival patterns of a given fixed length. We then study which of these models produce the qualitative phenomena we observe in real threads — particularly bimodality in the number of distinct commenters.

The first class of models \mathcal{F} we consider has the following basic structure for choosing who makes the j^{th} comment, reminiscent of the Chinese Restaurant Process [1]. With some fixed probability $p_j \geq 0$, we introduce a new participant; with probability $1 - p_j$ we select, according to some underlying probabilistic rule, a participant who has already appeared in the thread. (We refer to such participants as *re-entrants*). The re-entrant selection rule is assumed to be a randomized algorithm θ that takes a thread prefix as input and produces the name of an existing participant in the thread. This is a very general definition; depending on the choice of the function θ , we can define arbitrary rules, that, for example, pick a re-entrant uniformly, or according to “rich-get-richer” principles that favor people who have commented more in the past [15], or according to recency principles so that an individual’s selection probability decreases in the time since they last commented. Each model $\Omega(k, \theta, \mathbf{p})$ in this class \mathcal{F} is described by a thread length k , selection rule θ , and sequence of probabilities $\mathbf{p} = p_1, p_2, p_3, \dots, p_k$, $p_i \in (0, 1]$.

A Negative Result. Although \mathcal{F} initially seems reasonable and covers a large space, it turns out to be a poor fit to reality, because *none* of its members can yield the expansionary vs. focused bimodality that we found empirically in §3.

THEOREM 1. *Let $\Omega(k, \theta, \mathbf{p})$ be an arbitrary model in the class \mathcal{F} , and let X be a random variable equal to the number of distinct participants in a length- k thread t generated by $\Omega(k, \theta, \mathbf{p})$ (counting the initial poster). Then X has a unimodal density function: there is a number d^* such that $\Pr[X = d]$ is monotonically increasing for $d \leq d^*$ and monotonically decreasing for $d \geq d^*$.*

We omit the proof due to lack of space, but it consists essentially of projecting the arrival pattern onto a binary sequence that records only whether each participant is a re-entrant or not.

Models Exhibiting Bimodality. In view of this negative result, we seek an alternate class of models capable of generating arrival patterns that exhibit bimodality in the number of distinct commenters.

Arguably the simplest approach is to consider mixture models that have bimodality “built in”: We need only suppose that there are two distinct types of posts, one which concentrates the number of distinct participants on a small value, and the other which concentrates it on a large value, and that threads are constructed by

drawing one of the first type with fixed probability $\pi > 0$ or one of the second with probability $1 - \pi$.

While this mixture principle is presumably an important reason why we see bimodality in the real data, it is not the whole story. Indeed, we ran the following experiment to see whether the same type of post can lead both to focused and expansionary threads. As it turns out, the CNN link that was most shared among a large sample of Facebook users in the first quarter of 2012 was a report of Whitney Houston’s death. Although the set of threads spawned just by shares of this link is small by the standards of Figure 1, it is large in an absolute sense, and we observed in this controlled-content case the same sort of bimodality exhibited by threads overall: sometimes, the news provoked a series of “drive-by” comments when it was shared by a user, and other times, the same news prompted extended small-group discussion.

This finding motivates us to construct models of arrival patterns that produce the expansion/focus bimodality as a byproduct without assuming post type as its cause. To do this, we posit a type of internal symmetry-breaking during thread generation, taking inspiration from the theory of nonlinear urn processes [2]. In this new class of models, the probability that a new participant enters at step j depends on the identities of the participants in the first $j - 1$ steps. Intuitively, when there are many distinct participants, the process should make re-entry less likely, thereby producing momentum in the expansionary direction; when a few participants have each interacted multiple times, the process should make it harder for new participants to break in, thereby building up momentum in the conversational direction.

The class is parametrized by $\alpha \geq 1$ and $\beta \geq 0$. For each participant c already in the thread (including the original poster, 0), and each length $j \leq k$, each such existing participant will have a *weight* $w_j(c)$ after step j of the thread that controls their probability of providing the next comment. The fixed weight $\beta > 0$ controls the probability that a new participant arrives in the next step. We also impose the constraint that the same person never appears twice in a row.⁶

Generating the arrival pattern $\gamma = \gamma_1 \dots \gamma_k$ proceeds as follows. The first commenter will be labeled 1 (since we do not have the poster, labeled 0, provide the first comment too); so we initialize by setting $\gamma_1 = 1, w_1(0) = w_1(1) = 1$, and following this initialization we are positioned to determine the author of the second comment. In general, consider an arbitrary step $j < k$, and let c_j be the commenter in that step. We proceed as follows.

- (i) *Choose commenter $j + 1$.* We choose a participant (different from c_j) with probability proportional to the weights. Specifically: pre-existing participant $c \neq c_j$ is chosen with probability $w_j(c)/(\beta + \sum_{c' \neq c_j} w_j(c'))$, and a new participant is introduced into the thread with probability $\beta/(\beta + \sum_{c' \neq c_j} w_j(c'))$. We use c_{j+1} to denote the participant chosen for step $j + 1$.
- (ii) *Update weights.* If the participant c_{j+1} in step $j + 1$ is a re-entrant, we define $w_{j+1}(c_{j+1}) = \alpha w_j(c_{j+1})$, and leave all other weights unchanged. If instead c_{j+1} is new, we define $w_{j+1}(c_{j+1}) = 1$ and for all other pre-existing participants $c \neq c_{j+1}$ we reduce their weights by setting $w_{j+1}(c) = w_j(c)/\alpha$.

The key point is the weight update rule in part (ii). A new arrival suppresses the weight of all existing participants, making it less likely they will comment again and paving the way for further new

⁶This is essentially without loss of generality, since on the real threads we can also build a comparable representation where we collapse out consecutive occurrences of the same participant.




Figure 7: Density function of distinct participants in threads produced by our proposed family of processes with $\alpha = 1.5, 2$, and 4 ($\beta = 1$ and $k = 40$).

arrivals. On the other hand, when an existing participant provides the next comment, their increase in weight makes it more likely they will return, thereby promoting back-and-forth interaction.

We show via simulation that bimodality emerges naturally in this model. To paraphrase Langston Hughes, the number of distinct participants can dry up like a raisin in the sun, or it can explode. Figure 7 shows the empirical density function obtained through simulation for the number of distinct participants under multiple settings of the model parameters: we fix the length $k = 40$ and $\beta = 1$, and then we simulate the process with $\alpha = 1.5, 2$, and 4 . As we see there, bimodality emerges as α increases, which accords with intuition — larger values of α are more aggressive in amplifying both the focused and expansionary effects, and hence serve to bifurcate the process into its two modes more strongly.

The model appears to be quite challenging to analyze rigorously, and it is an interesting open question to prove that it produces bimodality, as well as to characterize the transition from unimodality to bimodality as we increase α . The model shares some properties with nonlinear urn processes [2], but also has ingredients that lie beyond what is usually needed for the analysis of such processes.

8. RELATED WORK

To our knowledge, there has not been prior consideration in the literature of the overall problem of algorithmic conversation curation — an emerging key component in enhancing user experience in current forms of on-line social interaction. This problem involves many issues, including those investigated in this paper: **(a)** determining which posts are interesting enough to bring to a user’s attention; **(b)** among discussions a user already has knowledge of, choosing which the user should continue to be updated about; and, indirectly, **(c)** understanding the structure of discussions, both to aid in the two issues just described and potentially for implications in user-interface design. Of course, there has been much valuable work on the first and last issue individually, which we now describe. (Our attention to **(b)** appears to be novel.) On **(a)**, we point out DeChoudhury et al.’s research [8] on the interestingness of Youtube comment threads, as measured by interestingness of topic and participants (not length), and Shmueli et al.’s work [8] on predicting which stories a particular user is most likely to comment on. Prior work on comment-volume prediction [3, 14, 23, 25, 26] is of course also quite relevant. How fast a piece of information spreads or diffuses [3, 4, 16, 18, 21] is another important aspect of interestingness. Quality of posts or comments, as determined by ratings, is potentially also relevant; see for example Siersdorfer et al. [22]. On **(c)**, there is intriguing work [12, 15] on structural characterizations of discussions when viewed as trees (not an approach we have taken in this paper, and arguably less natural as a model for discussions on sites like Facebook that have post-and-comment-interfaces). Different perspectives are taken by researchers look-




Figure 8: Considering only 8+-word posts that generated responses, for Facebook, the more distinctive the text of the original post, the more comments it garners, but for Wikipedia, which is more task-oriented, there is no such effect.

ing at characterizations of agreement and/or sentiment among comments [5, 11, 19, 20], and by sociological analyses of turn-taking in group conversations [10].

9. FUTURE WORK: DISTINCTIVENESS

We see many exciting future directions to pursue. Here, we briefly highlight two preliminary explorations into *distinctiveness features* that we believe hold promise, although we have not yet identified a way of applying them in their current form to improve prediction performance. The main idea is that the likelihood of a post’s text or early commenters should be informative.

Distinctiveness of Text. A basic property of a piece of text is its *likelihood* — whether its word choices look typical when compared to a reference collection, or whether its word choices are less likely and hence more *distinctive*. In recent work, measures of distinctiveness were shown to help in recognizing movie quotes that were deemed “memorable” in the sense of cultural penetration [6]. In our case, the question is the following: When a post contains unusual text, what should this lead us to estimate about the length of the resulting comment thread? There are intuitive arguments in both directions: some low-probability posts might generate discussion because they are provocative and unexpected, but others might simply be hard to understand and thus be mainly ignored.

We built a unigram language model from 3.5 million Facebook posts by authors whose posts weren’t in our main dataset; for each word w , the model provides a probability $p(w)$. We define a post’s *distinctiveness* to be the average over its tokens w of $\log(1/p(w))$; lower distinctiveness means a more likely post. Figure 8 shows the macro-averaged post length as a function of text distinctiveness, considering *only* posts containing at least 8 words⁷ and that received at least one comment in the case of Facebook or at least two for Wikipedia. For these particular subsets, our two Facebook populations exhibit a clear positive effect of the distinctiveness of the text, whereas for Wikipedia there seems to be no effect at all (which perhaps stems from the task-oriented nature of Wikipedian discussions). We note, however, that the effects become less clear if we include posts that turned out to generate no comments (or at most one on Wikipedia).

Distinctiveness of First Commenter. For a user u who posts regularly, a set of frequent commenters on u ’s threads often emerges — the people who generally weigh in when u says something. Thus, shifting from the likelihood of words to the likelihood of

⁷At 8 words and beyond, post distinctiveness becomes empirically almost independent of post length, disentangling the two features.




Figure 9: In Facebook, the more distinctive the first commenter is (in terms of not often being first to respond to the original poster), the longer the thread. Wikipedia is not depicted due to sparseness, but the overall trend is the opposite.

users, it makes sense to ask about the effect on thread length of the first commenter’s distinctiveness — the extent to which this commenter is usually or rarely first in one of u ’s threads. Again, there are intuitive arguments each way: if the first commenter v is someone who’s often a first commenter on u ’s posts, then v is presumably familiar to both u and the audience for u ’s comments, which could make it easier for the thread to grow; but it may also be socially easier to let v ’s comment pass by without much activity.

In Figure 9 we show for Facebook the expected thread length as a function of the fraction of times the first commenter was *not* the first to respond to the original poster’s posts⁸. We see a clear upward trend: when someone you rarely hear from first is in fact the first to comment on your post, on average it foreshadows a longer thread, perhaps because this indicates that the post has greater reach.

We note that Wikipedia appears to exhibit the opposite behavior. We do not depict the Wikipedia results due to sparseness of recurring first commenters, but when restricting to users with at least 10 posts and binning the distinctiveness values, we see a significant decreasing trend.

10. CONCLUSIONS

Motivated by the growing role of automated mechanisms to manage users’ interactions with on-line discussions, we have identified and studied two key problems in the curation of such discussions. The first of these, length prediction, is related to earlier studies of comment volume on blog and news sites, but it acquires additional complexity in our context due to the heterogeneity we find in long threads, which can either be focused on a few participants or expand to reach many. The second problem, re-entry prediction, has to our knowledge not been formulated previously; it is a crucial issue in applications that must decide when to notify users about updates to discussions in which they have participated.

We see these two problems as helping to define the contours of the problem of conversational curation more broadly, and as such the results here suggest a range of further open questions. Among these are a deeper understanding of the features that can help predict the trajectory of an on-line discussion from its early stages, and the integration of these techniques into systems that deliver discussion-oriented content to users in on-line applications.

Acknowledgments We thank Tom Lento and the anonymous reviewers for their helpful comments. Supported in part by NSF grant IIS-0910664. This work was done while the last author was at Cornell University.

⁸For the uniform population, the plot only consider users with at least ten posts, although different threshold values do not greatly affect the resulting trends.

References

- [1] D. J. Aldous. Exchangeability and related topics. *École d’Été de Probabilités de Saint-Flour XIII (Lec. Notes. Math. vol 1117)*, 1985.
- [2] W. B. Arthur. Competing technologies, increasing returns, and lock-in by historical events. *The Economic Journal*, 1989.
- [3] Y. Artzi, P. Pantel, M. Gamon. Predicting responses to microblog posts. *NAACL (short paper)*, 2012.
- [4] E. Bakshy, J. M. Hofman, W. A. Mason, D. J. Watts. Everyone’s an influencer: Quantifying influence on Twitter. *WSDM*, 2011.
- [5] A. Chmiel, J. Sienkiewicz, M. Thelwall, G. Paltoglou, K. Buckley, A. Kappas, J. A. Holyst. Collective emotions online and their influence on community life. *PLoS One*, 2011.
- [6] C. Danescu-Niculescu-Mizil, J. Cheng, J. Kleinberg, L. Lee. You had me at hello: How phrasing affects memorability. *ACL*, 2012.
- [7] C. Danescu-Niculescu-Mizil, L. Lee, B. Pang, J. Kleinberg. Echoes of power: Language effects and power differences in social interaction. *WWW*, 2012.
- [8] M. De Choudhury, H. Sundaram, A. John, D. D. Seligmann. What makes conversations interesting?: Themes, participants and consequences of conversations in online social media. *WWW*, 2009.
- [9] J. Friedman, T. Hastie, R. Tibshirani. Regularization paths for generalized linear models via coordinate descent. *J. Statistical Software*, 2010.
- [10] D. R. Gibson. Marking the turn: Obligation, engagement, and alienation in group discussions. *Social Psychology Quarterly*, 2010.
- [11] E. Gilbert, T. Bergstrom, K. Karahalios. Blogs are Echo Chambers: Blogs are Echo Chambers. *HICSS*, 2009.
- [12] V. Gómez, A. Kaltenbrunner, V. López. Statistical analysis of the social network and discussion threads in Slashdot. *WWW*, 2008.
- [13] V. Gómez, H. J. Kappen, A. Kaltenbrunner. Modeling the structure and evolution of discussion cascades. *Hypertext*, 2011.
- [14] M. Guerini, C. Strapparava, G. Özal. Exploring text virality in social networks. *ICWSM (poster)*, 2011.
- [15] R. Kumar, M. Mahdian, M. McGlohon. Dynamics of conversations. *KDD*, 2010.
- [16] H. Kwak, C. Lee, H. Park, S. Moon. What is Twitter, a social network or a news media? *WWW*, 2010.
- [17] D. Laniado, R. Tasso, Y. Volkovich, A. Kaltenbrunner. When the Wikipedians talk: Network and tree structure of Wikipedia discussion pages. *ICWSM*, 2011.
- [18] K. Lerman, R. Ghosh. Information contagion: An empirical study of the spread of news on Digg and Twitter social networks. *ICWSM*, 2010.
- [19] G. Mishne, N. Glance. Leave a reply: An analysis of weblog comments. *Wksp. on the Weblogging ecosystem*, 2006.
- [20] S. Park, M. Ko, J. Kim, Y. Liu, J. Song. The politics of comments: predicting political orientation of news stories with commenters’ sentiment patterns. *CSCW*, 2011.
- [21] D. M. Romero, B. Meeder, J. Kleinberg. Differences in the mechanics of information diffusion across topics: Idioms, political hashtags, and complex contagion on Twitter. *WWW*, 2011.
- [22] S. Siersdorfer, S. Chelaru, W. Nejdl, J. San Pedro. How useful are your comments?: Analyzing and predicting YouTube comments and comment ratings. *WWW*, 2010.
- [23] M. Tsagkias, W. Weerkamp, M. Rijke. Predicting the volume of comments on online news stories. *CIKM*, 2009.
- [24] J. Ugander, L. Backstrom, C. Marlow, J. Kleinberg. Structural diversity in social contagion. *PNAS*, 2012.
- [25] C. Wang, M. Ye, B. A. Huberman. From user comments to on-line conversations. *KDD*, 2012.
- [26] T. Yano, N. A. Smith. What’s worthy of comment? Content and comment volume in political blogs. *ICWSM*, 2010.